

Neural Network For Letter Recoignition From Scratch

December 28, 2022

```
[1]: import numpy as np
import pandas as pd
from matplotlib import figure
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import OneHotEncoder
```

```
[2]: df= pd.read_csv('letter-recognition.data', header=None)
print(df.shape)
df.head()
```

(20000, 17)

```
[2]:  0   1   2   3   4   5   6   7   8   9  10  11  12  13  14  15  16
0   T   2   8   3   5   1   8  13   0   6   6  10   8   0   8   0   8
1   I   5  12   3   7   2  10   5   5   4  13   3   9   2   8   4  10
2   D   4  11   6   8   6  10   6   2   6  10   3   7   3   7   3   9
3   N   7  11   6   6   3   5   9   4   6   4   4  10   6  10   2   8
4   G   2   1   3   1   1   8   6   6   6   6   5   9   1   7   5  10
```

```
[3]: df.describe()
```

```
[3]:
```

	1	2	3	4	5 \
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	4.023550	7.035500	5.121850	5.37245	3.505850
std	1.913212	3.304555	2.014573	2.26139	2.190458
min	0.000000	0.000000	0.000000	0.00000	0.000000
25%	3.000000	5.000000	4.000000	4.00000	2.000000
50%	4.000000	7.000000	5.000000	6.00000	3.000000
75%	5.000000	9.000000	6.000000	7.00000	5.000000
max	15.000000	15.000000	15.000000	15.00000	15.000000

	6	7	8	9	10 \
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	6.897600	7.500450	4.628600	5.178650	8.282050
std	2.026035	2.325354	2.699968	2.380823	2.488475
min	0.000000	0.000000	0.000000	0.000000	0.000000

25%	6.000000	6.000000	3.000000	4.000000	7.000000
50%	7.000000	7.000000	4.000000	5.000000	8.000000
75%	8.000000	9.000000	6.000000	7.000000	10.000000
max	15.000000	15.000000	15.000000	15.000000	15.000000

	11	12	13	14	15 \
count	20000.000000	20000.000000	20000.000000	20000.000000	20000.000000
mean	6.45400	7.929000	3.046100	8.338850	3.691750
std	2.63107	2.080619	2.332541	1.546722	2.567073
min	0.00000	0.000000	0.000000	0.000000	0.000000
25%	5.00000	7.000000	1.000000	8.000000	2.000000
50%	6.00000	8.000000	3.000000	8.000000	3.000000
75%	8.00000	9.000000	4.000000	9.000000	5.000000
max	15.00000	15.000000	15.000000	15.000000	15.000000

	16
count	20000.000000
mean	7.80120
std	1.61747
min	0.00000
25%	7.00000
50%	8.00000
75%	9.00000
max	15.00000

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 17 columns):
#   Column  Non-Null Count  Dtype
---  -
0   0        20000 non-null     object
1   1        20000 non-null     int64
2   2        20000 non-null     int64
3   3        20000 non-null     int64
4   4        20000 non-null     int64
5   5        20000 non-null     int64
6   6        20000 non-null     int64
7   7        20000 non-null     int64
8   8        20000 non-null     int64
9   9        20000 non-null     int64
10  10       20000 non-null     int64
11  11       20000 non-null     int64
12  12       20000 non-null     int64
13  13       20000 non-null     int64
14  14       20000 non-null     int64
15  15       20000 non-null     int64
```

```

16 16      20000 non-null int64
dtypes: int64(16), object(1)
memory usage: 2.6+ MB

```

```

[5]: # label encoding the class labels
from sklearn.preprocessing import LabelEncoder
label_encoder= LabelEncoder()
df[0]= label_encoder.fit_transform(df[0])
print(df[0].unique())
df.head()

```

```

[19  8  3 13  6 18  1  0  9 12 23 14 17  5  2  7 22 11 15  4 21 24 16 20
 10 25]

```

```

[5]:
   0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16
0 19  2  8  3  5  1  8 13  0  6  6 10  8  0  8  0  8
1  8  5 12  3  7  2 10  5  5  4 13  3  9  2  8  4 10
2  3  4 11  6  8  6 10  6  2  6 10  3  7  3  7  3  9
3 13  7 11  6  6  3  5  9  4  6  4  4 10  6 10  2  8
4  6  2  1  3  1  1  8  6  6  6  6  5  9  1  7  5 10

```

```

[6]: df_test= df[[0]]
print(df_test.shape)
df_test.head()

```

```

(20000, 1)

```

```

[6]:
   0
0 19
1  8
2  3
3 13
4  6

```

```

[7]: df_train= df.drop(0, axis=1, inplace=False)
print(df_train.shape)
df_train.head()

```

```

(20000, 16)

```

```

[7]:
   1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16
0  2  8  3  5  1  8 13  0  6  6 10  8  0  8  0  8
1  5 12  3  7  2 10  5  5  4 13  3  9  2  8  4 10
2  4 11  6  8  6 10  6  2  6 10  3  7  3  7  3  9
3  7 11  6  6  3  5  9  4  6  4  4 10  6 10  2  8
4  2  1  3  1  1  8  6  6  6  6  5  9  1  7  5 10

```

```

[8]: from sklearn.model_selection import train_test_split

```

```
[9]: X_train, X_rem, y_train, y_rem = train_test_split(df_train, df_test,
    ↪train_size=0.8)

X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)

print("x_train dataset size: ",X_train.shape)
print("y_train dataset size: ",y_train.shape)
print("x_validation dataset size: ",X_valid.shape)
print("y_validation dataset size: ",y_valid.shape)
print("x_test dataset size: ",X_test.shape)
print("y_test dataset size: ",y_test.shape)
```

```
x_train dataset size: (16000, 16)
y_train dataset size: (16000, 1)
x_validation dataset size: (2000, 16)
y_validation dataset size: (2000, 1)
x_test dataset size: (2000, 16)
y_test dataset size: (2000, 1)
```

```
[10]: print("number of unique classes in the dataset is- ", len(df[0].unique()))
```

```
number of unique classes in the dataset is- 26
```

```
[11]: df[0].value_counts()
```

```
[11]: 20      813
      3      805
      15     803
      19     796
      12     792
      0     789
      23     787
      24     786
      13     783
      16     783
      5     775
      6     773
      4     768
      1     766
      21     764
      11     761
      17     758
      8     755
      14     753
      22     752
      18     748
      9     747
      10     739
```

```


2      736
7      734
25     734
Name: 0, dtype: int64

```

```

[12]: plt.figure(figsize = (13,6))
plt.title("COUNT OF LETTERS")
sns.countplot(x=0, data= df)
plt.show()

```



```

[13]: # we can observe that the dataset is balanced and dont have any null values so we
      ↪ proceed with ot

```

1 Building the NN

```

[14]: # it has 2 layers i.e. 1 input, 1 hidden and 1 output layer
      # lets assume
      # 40 neurons in the first hidden layer
      # output layer is a 26 class softmax layer for multiclass classification

```

weight matrix between input layer and 1st hidden layer has size $[16 \times 40]$ bias vector between input layer and 1st hidden layer is of size $[40 \times 1]$

weight matrix between 1st hidden layer and output softmax layer is $[40 \times 26]$ bias vector between 1st hidden layer and output softmax layer is of size $[26 \times 1]$

```

[15]: # converting all the datasets into arrays
x_train_arr = np.array(X_train)

```

```
y_train_arr = np.array(y_train)
x_test_arr = np.array(X_test)
y_test_arr = np.array(y_test)
x_valid_arr = np.array(X_valid)
y_valid_arr = np.array(y_valid)
```

```
[16]: onehot= OneHotEncoder()
```

```
[17]: y_train_arr = onehot.fit_transform(y_train_arr.reshape(-1,1))
y_train_arr = y_train_arr.toarray()
y_train_arr
```

```
[17]: array([[0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.]])
```

```
[18]: y_test_arr = onehot.fit_transform(y_test_arr.reshape(-1,1))
y_test_arr = y_test_arr.toarray()
y_train_arr
```

```
[18]: array([[0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.]])
```

```
[19]: y_valid_arr = onehot.fit_transform(y_valid_arr.reshape(-1,1))
y_valid_arr = y_valid_arr.toarray()
y_valid_arr
```

```
[19]: array([[0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 1.],
          [0., 0., 0., ..., 0., 0., 0.],
          ...,
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.],
          [0., 0., 0., ..., 0., 0., 0.]])
```

```
[20]: class activation_function:
      def __init__(self):
          self.name = None
          self.output = None
```

```

def tanh(self, input_data):
    return np.tanh(input_data)

def grad_tanh(self, input_data):
    return 1-np.tanh(input_data)**2

def relu(self, x):
    self.output = np.maximum(0,x)
    return self.output

def grad_relu(self, x):
    self.output[self.output>0] = 1
    return self.output

def activation(self, name, x):
    self.name = name
    if self.name == "relu":
        return self.relu(x)
    elif self.name == "tanh":
        return self.tanh(x)

def grad_activation(self,x):
    if self.name == "relu":
        return self.grad_relu(x)
    elif self.name == "tanh":
        return self.grad_tanh(x)

```

[21]: *# defining all three activation functions*

```

def activation_relu(x):
    return np.maximum(0,x)

def activation_tanh(x):
    return np.tanh(x)

# def activation_softmax(x):
#     exp_values= np.exp(x - np.max(x, axis=1, keepdims=True))
#     probabilities= exp_values/ np.sum(exp_values, axis= 1, keepdims=True)
#     return probabilities

```

[22]: *#defining the gradients of both the activation functions for updating the weights*

```

def gradient_relu(x):
    x[x>0]=1
    return x

```

```
def gradient_tanh(x):
    return (1- np.tanh(x)**2)

# def gradient_softmax(x):
#     jacobian= np.diag(x)
#     for i in range(len(jacobian)):
#         for j in range(len(jacobian)):
#             if i == j:
#                 jacobian[i][j] = x[i] * (1-x[i])
#             else:
#                 jacobian[i][j] = -x[i]*x[j]
#     return jacobian
```

```
[23]: # #defining the adam optimizer for faster convergence

# def optimizer_adam(weight, alpha=0.001, beta1= 0.9, beta2= 0.999, epsilon=
→1e-8):
#     t=0

#     m_t=0
#     v_t=0

#     m_t_hat= m_t
#     v_t_hat= v_t

#     num_steps=1000

#     while t in range(1,num_steps):
#         dw= compute_gradient(weight)

#         m_t= beta1*m_t+ (1-beta1)*dw
#         v_t= beta2*v_t+ (1-beta2)*dw*dw

#         m_t_hat= m_t/(1- beta1**t)
#         v_t_hat= v_t/(1-beta2**t)

#         weight= weight- alpha*(m_t_hat/(v_t_hat.sqrt()+ epsilon))

#     return weight
```

```
[24]: #defining a layer

class layer:
    def __init__(self):
        self.input = None
        self.output = None
```



```

# computes the output Y of a layer for a given input X
def forward_propagation(self, input):
    pass

# computes dE/dX for a given dE/dY (and update parameters if any)
def backward_propagation(self, output_error, learning_rate):
    pass

```

[25]: # the fully connected layer

```

class fully_connected_layer:
    def __init__(self, input_size, output_size):
        # xavier weight initialization
        # input_size = fan_in
        # output_size = fan_out
        self.weights = np.random.randn(input_size, output_size) / np.
        sqrt(input_size)
        self.bias = np.random.randn(1, output_size)

    # returns output for a given input
    def forward_propagation(self, input_data):
        self.input = input_data.reshape(1, -1)
        self.output = np.dot(self.input, self.weights) + self.bias
        return self.output

    # computes dE/dW, dE/dB for a given output_error=dE/dY. Returns
    input_error=dE/dX.
    def backward_propagation(self, output_gradient, learning_rate, optimizer, T):
        input_gradient = np.dot(output_gradient, self.weights.T)
        weights_gradient = np.dot(self.input.T, output_gradient)
        bias_gradient = output_gradient

        # updating the parameters
        self.weights = self.weights - (learning_rate * weights_gradient)
        self.bias = self.bias - (learning_rate * output_gradient)

        if optimizer=='adam':
            opt = AdamOptim(eta=learning_rate, beta1=0.9, beta2=0.999,
            epsilon=1e-8)
            self.weights, self.bias = opt.update(t=T, w = self.weights, b = self.
            bias, dw = weights_gradient, db = output_gradient)

        return input_gradient

```

[26]: # # the actiation layer

```

# class activation_layer:

```

```

#     def __init__(self, activation, grad_activation):
#         self.activation = activation
#         self.grad_activation = grad_activation

#     # returns the activated input
#     def forward_propagation(self, input_data):
#         self.input = input_data
#         self.output = self.activation(self.input)
#         return self.output

#     # Returns input_error=dE/dX for a given output_error=dE/dY.
#     # learning_rate is not used because there is no "learnable" parameters.
#     def backward_propagation(self, output_gradient, learning_rate, optimizer, T):
#         return self.grad_activation(self.input) * output_gradient

```

```

[27]: # the activation layer

class activation_layer:
    def __init__(self, act_func):
        self.input = None
        self.output = None
        self.act_func = act_func
        self.act = activation_function()

    def forward_propagation(self, input_data):
        self.input = input_data
        self.output = self.act.activation(self.act_func, input_data)
        return self.output

    def backward_propagation(self, output_gradient, learning_rate, optimizer, T):
        return self.act.grad_activation(self.input)*output_gradient

```

```

[28]: # defining the softmax activation layer

class softmax_layer:
    def __init__(self, input_size):
        self.input = None
        self.output = None
        self.input_size = input_size

    def forward_propagation(self, input_data):
        self.input = input_data
        exp = np.exp(self.input - np.max(self.input, axis=1, keepdims=True))
        exp_sum = np.sum(exp, axis=1, keepdims=True)
        self.output = exp/(exp_sum)
        return self.output

```

```

def backward_propagation(self, output_data, learning_rate, optimizer, T):
    dz_x = np.zeros((self.input_size,self.input_size))
    for i in range(self.input_size):
        for j in range(self.input_size):
            if i == j:
                dz_x[i,j] = self.output[0,i]*(1-self.output[0,i])
            else:
                dz_x[i,j] = -self.output[0,i]*self.output[0,j]

    dl_x = np.dot(output_data,dz_x)
    return dl_x

```

```

[29]: # defining the loss function
def cross_entropy_loss(y_true, y_pred):
    loss = np.sum(-np.log(y_pred+1e-20)*y_true)
    return loss

def grad_cross_entropy_loss(y_true, y_pred):
    grad_loss = -y_true/(y_pred + 1e-20)
    return grad_loss

```

```

[30]: # defining the adam optimizer
class AdamOptim():
    def __init__(self, eta=0.01, beta1=0.9, beta2=0.999, epsilon=1e-8):
        self.m_dw, self.v_dw = 0, 0
        self.m_db, self.v_db = 0, 0
        self.beta1 = beta1
        self.beta2 = beta2
        self.epsilon = epsilon
        self.eta = eta

    def update(self, t, w, b, dw, db):
        ## dw, db are from current minibatch
        ## momentum beta 1
        # *** weights *** #
        self.m_dw = self.beta1*self.m_dw + (1-self.beta1)*dw
        # *** biases *** #
        self.m_db = self.beta1*self.m_db + (1-self.beta1)*db

        ## rms beta 2
        # *** weights *** #
        self.v_dw = self.beta2*self.v_dw + (1-self.beta2)*(dw**2)
        # *** biases *** #
        self.v_db = self.beta2*self.v_db + (1-self.beta2)*(db**2)

        ## bias correction
        m_dw_corr = self.m_dw/(1-self.beta1**t)
        m_db_corr = self.m_db/(1-self.beta1**t)

```

```

v_dw_corr = self.v_dw/(1-self.beta2**t)
v_db_corr = self.v_db/(1-self.beta2**t)
#     print(v_db_corr)

## update weights and biases
w = w - self.eta*(m_dw_corr/(np.sqrt(v_dw_corr)+self.epsilon))
b = b - self.eta*(m_db_corr/(np.sqrt(v_db_corr)+self.epsilon))
return w, b

```

```

[31]: class MLP:
    def __init__(self):
        self.layers = []
        self.loss = None
        self.grad_loss= None
        self.optimizer = None
        self.train_loss = []
        self.validation_loss = []
        self.train_accuracy= []
        self.validation_accuracy = []

        # add layer to network
    def add(self, layer):
        self.layers.append(layer)

        # set loss to use
    def use(self, loss, grad_loss):
        self.loss = loss
        self.grad_loss= grad_loss

        # predict output for given input
    def predict(self, input_data):
        # sample dimension first
        samples = len(input_data)
        result = []

        # run network over all samples
        for i in range(samples):
            # forward propagation
            output = input_data[i]
            for layer in self.layers:
                output = layer.forward_propagation(output)
            output= (output==np.max(output))*1
            result.append(output[0])

        return result

    # train the network

```

```

def fit(self, x_train, y_train, x_valid, y_valid, epochs, learning_rate,
optimizer="adam"):
    # sample dimension first
    samples = len(x_train)
    self.optimizer= optimizer
    # training loop
    for i in range(epochs):
        err = 0
        for j in range(samples):
            # forward propagation
            output = x_train[j]
            for layer in self.layers:
                output = layer.forward_propagation(output)

            # compute loss
            err += self.loss(y_train[j], output)

            # backward propagation
            error = self.grad_loss(y_train[j], output)
            for layer in reversed(self.layers):
                error = layer.backward_propagation(error, learning_rate,
optimizer= self.optimizer, T=i+1)

        # training loss
        err /= samples
        self.train_loss.append(err)
        print("epoch- {}/{}", loss- {}".format(i,epochs,err))

        # accuracy for train set
        y_pred = self.predict(x_train)
        y_pred = onehot.inverse_transform(y_pred)
        y_pred = y_pred.flatten()

        y_true = y_train
        y_true = onehot.inverse_transform(y_true)
        y_true = y_true.flatten()

        self.train_accuracy.append(accuracy_score(y_pred, y_true))

    #Validation loss
    valid_samples= len(x_valid)
    err_valid= 0
    for j in range(valid_samples):
        output = x_valid[j]
        for layer in self.layers:
            output = layer.forward_propagation(output)
        err_valid += self.loss(y_valid[j], output)

```

```

err_valid /= valid_samples
self.validation_loss.append(err_valid)

# accuracy for validation set
y_valid_pred = self.predict(x_valid)
y_valid_pred = onehot.inverse_transform(y_valid_pred)
y_valid_pred = y_valid_pred.flatten()

y_valid_true = y_valid
y_valid_true = onehot.inverse_transform(y_valid_true)
y_valid_true = y_valid_true.flatten()

self.validation_accuracy.append(accuracy_score(y_valid_pred,
↪y_valid_true))

def plot(self):
    # plotting the graphs for loss vs epochs
    plt.figure(figsize=(15,8))
    plt.subplot(1,2,1)
    plt.plot(self.train_loss, label = "training set")
    plt.plot(self.validation_loss, label = "validation set")
    plt.title("loss vs epochs")
    plt.xlabel("epochs")
    plt.ylabel("loss")
    plt.legend()

    # plotting the graphs for accuracy vs epochs
    plt.subplot(1,2,2)
    plt.plot(self.train_accuracy, label = "training set")
    plt.plot(self.validation_accuracy, label = "validation set")
    plt.title("accuracy vs epochs")
    plt.xlabel("epochs")
    plt.ylabel("accuracy")
    plt.legend()

```

2 for tanh activation

```

[32]: # neural network for tanh activation function
nn= MLP()
nn.add(fully_connected_layer(16, 40))
nn.add(activation_layer(act_func="tanh"))
#nn.add(activation_layer(activation_tanh, gradient_tanh))
nn.add(fully_connected_layer(40, 26))
nn.add(softmax_layer(26))
nn.use(cross_entropy_loss, grad_cross_entropy_loss)

```

```

nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=20,
      ↪learning_rate=0.001, optimizer="adam")
nn.plot()


y_test_pred = np.array(nn.predict(x_test_arr))
y_test_pred = onehot.inverse_transform(y_test_pred)
y_test_true = onehot.inverse_transform(y_test_arr)
print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))

```

```

epoch- 0/20, loss- 2.4848552407795155
epoch- 1/20, loss- 1.770108385888565
epoch- 2/20, loss- 1.651964738313998
epoch- 3/20, loss- 1.6159811464985474
epoch- 4/20, loss- 1.5928636730328514
epoch- 5/20, loss- 1.5596862088299344
epoch- 6/20, loss- 1.5154159481811098
epoch- 7/20, loss- 1.4750465531988715
epoch- 8/20, loss- 1.460720551326636
epoch- 9/20, loss- 1.4594251237406901
epoch- 10/20, loss- 1.4455482761203344
epoch- 11/20, loss- 1.4556119168477015
epoch- 12/20, loss- 1.4458132059837878
epoch- 13/20, loss- 1.4498966800311928
epoch- 14/20, loss- 1.4238306677000039
epoch- 15/20, loss- 1.4039495491951772
epoch- 16/20, loss- 1.3935314119991875
epoch- 17/20, loss- 1.3880206577278735
epoch- 18/20, loss- 1.404275128613891
epoch- 19/20, loss- 1.4062536477525163
Accuracy on test set: 0.674

```



3 for relu activation

```
[37]: # neural network for relu activation function
nn= MLP()
nn.add(fully_connected_layer(16, 40))
nn.add(activation_layer(act_func="relu"))
#nn.add(activation_layer(activation_relu, gradient_relu))
nn.add(fully_connected_layer(40, 26))
nn.add(softmax_layer(26))
nn.use(cross_entropy_loss, grad_cross_entropy_loss)
nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=20,
      ↪learning_rate=1e-4, optimizer="adam")
nn.plot()

y_test_pred = np.array(nn.predict(x_test_arr))
y_test_pred = onehot.inverse_transform(y_test_pred)
y_test_true = onehot.inverse_transform(y_test_arr)
print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))
```


```
epoch- 0/20, loss- 2.800645124648064
epoch- 1/20, loss- 1.9573487388589554
epoch- 2/20, loss- 1.6613207721119385
epoch- 3/20, loss- 1.5006014524197293
epoch- 4/20, loss- 1.4020208037404536
epoch- 5/20, loss- 1.3406659913878556
epoch- 6/20, loss- 1.3021977272718872
```



```

epoch- 7/20, loss- 1.2718083890888852
epoch- 8/20, loss- 1.2511874165531032
epoch- 9/20, loss- 1.2357377567541372
epoch- 10/20, loss- 1.2255004479148295
epoch- 11/20, loss- 1.2193503086813335
epoch- 12/20, loss- 1.2158565102584344
epoch- 13/20, loss- 1.2139335391813948
epoch- 14/20, loss- 1.2149939346970002
epoch- 15/20, loss- 1.2173295336687786
epoch- 16/20, loss- 1.220953492865633
epoch- 17/20, loss- 1.225315314281561
epoch- 18/20, loss- 1.2328576003504141
epoch- 19/20, loss- 1.2439312460238066
Accuracy on test set: 0.635

```



4 for different learning rates on tanh activation

```

[34]: lr=[1e-5,1e-4,0.001,0.01,0.1]
for i in lr:
    print("training for learning rate",i)
    print("-----")
    nn= MLP()
    nn.add(fully_connected_layer(16, 40))
    nn.add(activation_layer(act_func="tanh"))
    #nn.add(activation_layer(activation_relu, gradient_relu))
    nn.add(fully_connected_layer(40, 26))

```

```

nn.add(softmax_layer(26))
nn.use(cross_entropy_loss, grad_cross_entropy_loss)
nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=20,
→learning_rate=i, optimizer="adam")
nn.plot()

y_test_pred = np.array(nn.predict(x_test_arr))
y_test_pred = onehot.inverse_transform(y_test_pred)
y_test_true = onehot.inverse_transform(y_test_arr)
print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))
print("\n")

```

training for learning rate 1e-05

```

-----
epoch- 0/20, loss- 3.6507989615111143
epoch- 1/20, loss- 3.3902791159594217
epoch- 2/20, loss- 3.2832223832479825
epoch- 3/20, loss- 3.2134975221111484
epoch- 4/20, loss- 3.163537802493247
epoch- 5/20, loss- 3.1241895436310485
epoch- 6/20, loss- 3.0905104451033045
epoch- 7/20, loss- 3.060747430640881
epoch- 8/20, loss- 3.0330097263403415
epoch- 9/20, loss- 3.0065955798279407
epoch- 10/20, loss- 2.981069146093323
epoch- 11/20, loss- 2.956338682728828
epoch- 12/20, loss- 2.931674579740753
epoch- 13/20, loss- 2.9061196235062248
epoch- 14/20, loss- 2.878782635976009
epoch- 15/20, loss- 2.8511179212017197
epoch- 16/20, loss- 2.82498457217815
epoch- 17/20, loss- 2.7995995148710726
epoch- 18/20, loss- 2.774438691028938
epoch- 19/20, loss- 2.749230885574704
Accuracy on test set: 0.215

```

training for learning rate 0.0001

```

-----
epoch- 0/20, loss- 3.2320028393054367
epoch- 1/20, loss- 2.9010764373407216
epoch- 2/20, loss- 2.6771952234387153
epoch- 3/20, loss- 2.5319819371778136
epoch- 4/20, loss- 2.3966475123225788
epoch- 5/20, loss- 2.274682416329273
epoch- 6/20, loss- 2.1658703533141273
epoch- 7/20, loss- 2.0747467239867925
epoch- 8/20, loss- 1.996132584750352

```

```
epoch- 9/20, loss- 1.9215871100361723
epoch- 10/20, loss- 1.8549408065297421
epoch- 11/20, loss- 1.8024457302338028
epoch- 12/20, loss- 1.760518797588856
epoch- 13/20, loss- 1.7253433051144413
epoch- 14/20, loss- 1.694341243537717
epoch- 15/20, loss- 1.666307184695178
epoch- 16/20, loss- 1.6414240956341197
epoch- 17/20, loss- 1.6191945131940597
epoch- 18/20, loss- 1.5993962889691378
epoch- 19/20, loss- 1.581512127587847
Accuracy on test set: 0.5445
```

training for learning rate 0.001

```
-----
epoch- 0/20, loss- 2.585573293003264
epoch- 1/20, loss- 1.831271402698168
epoch- 2/20, loss- 1.6137217545762945
epoch- 3/20, loss- 1.476979410185661
epoch- 4/20, loss- 1.4205451408626988
epoch- 5/20, loss- 1.3956305428458786
epoch- 6/20, loss- 1.3729319141755978
epoch- 7/20, loss- 1.3499768465582915
epoch- 8/20, loss- 1.3486021684000515
epoch- 9/20, loss- 1.3419190794052553
epoch- 10/20, loss- 1.314436708213282
epoch- 11/20, loss- 1.3047411350332974
epoch- 12/20, loss- 1.3153285458627295
epoch- 13/20, loss- 1.3143125338214823
epoch- 14/20, loss- 1.2971461882859967
epoch- 15/20, loss- 1.290838347161119
epoch- 16/20, loss- 1.2956511500705248
epoch- 17/20, loss- 1.2849351944787162
epoch- 18/20, loss- 1.2914247332617703
epoch- 19/20, loss- 1.2887430021721815
Accuracy on test set: 0.6865
```





training for learning rate 0.01





```
-----
epoch- 0/20, loss- 2.800153680157735
epoch- 1/20, loss- 2.5325232409537493
epoch- 2/20, loss- 2.3808311538502207
epoch- 3/20, loss- 2.2051589141338286
epoch- 4/20, loss- 2.2043700291283965
epoch- 5/20, loss- 2.211845481931634
epoch- 6/20, loss- 2.102422320430022
```


```
epoch- 7/20, loss- 2.0160011820139863
epoch- 8/20, loss- 2.0668799162578524
epoch- 9/20, loss- 1.9675458424502603
epoch- 10/20, loss- 1.9600399723680562
epoch- 11/20, loss- 1.9347114269485588
epoch- 12/20, loss- 1.9073709163780088
epoch- 13/20, loss- 1.8741254945618087
epoch- 14/20, loss- 1.849342603097563
epoch- 15/20, loss- 1.859689923432603
epoch- 16/20, loss- 1.8406772534425215
epoch- 17/20, loss- 1.8327936050019198
epoch- 18/20, loss- 1.9019239308223903
epoch- 19/20, loss- 1.8641683920966312
Accuracy on test set: 0.5815
```

training for learning rate 0.1

```
-----
epoch- 0/20, loss- 12.090184514752083
epoch- 1/20, loss- 10.499906073962102
epoch- 2/20, loss- 9.814617284393124
epoch- 3/20, loss- 9.405926745474895
epoch- 4/20, loss- 9.138706091613674
epoch- 5/20, loss- 8.955067089495921
epoch- 6/20, loss- 8.827011550890841
epoch- 7/20, loss- 8.737226127144089
epoch- 8/20, loss- 8.67645939920827
epoch- 9/20, loss- 8.636320598248565
epoch- 10/20, loss- 8.612367957695774
epoch- 11/20, loss- 8.601083653461352
epoch- 12/20, loss- 8.599667422579149
epoch- 13/20, loss- 8.606380383082977
epoch- 14/20, loss- 8.620142881410128
epoch- 15/20, loss- 8.638903191075485
epoch- 16/20, loss- 8.66203903653551
epoch- 17/20, loss- 8.688765201291044
epoch- 18/20, loss- 8.718437929029783
epoch- 19/20, loss- 8.750524642678858
Accuracy on test set: 0.042
```







5 for different learning rates on relu activation

```
[35]: lr=[1e-5,1e-4,0.001,0.01,0.1]
for i in lr:
    print("training for learning rate",i)
    print("-----")
    nn= MLP()
    nn.add(fully_connected_layer(16, 40))
    nn.add(activation_layer(act_func="relu"))
    #nn.add(activation_layer(activation_relu, gradient_relu))
    nn.add(fully_connected_layer(40, 26))
    nn.add(softmax_layer(26))
    nn.use(cross_entropy_loss, grad_cross_entropy_loss)
    nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=20,
    ↪learning_rate=i, optimizer="adam")
    nn.plot()

    y_test_pred = np.array(nn.predict(x_test_arr))
    y_test_pred = onehot.inverse_transform(y_test_pred)
    y_test_true = onehot.inverse_transform(y_test_arr)
    print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))
    print("\n")
```

training for learning rate 1e-05

epoch- 0/20, loss- 4.664039274644299

epoch- 1/20, loss- 3.540341895821613
epoch- 2/20, loss- 3.256182359364265
epoch- 3/20, loss- 3.0874346285416263
epoch- 4/20, loss- 2.954579316064428
epoch- 5/20, loss- 2.8376539498696904
epoch- 6/20, loss- 2.731186502791021
epoch- 7/20, loss- 2.6336030261319117
epoch- 8/20, loss- 2.543621969032915
epoch- 9/20, loss- 2.460210536802279
epoch- 10/20, loss- 2.383183119098905
epoch- 11/20, loss- 2.3114287125939277
epoch- 12/20, loss- 2.2448279374174707
epoch- 13/20, loss- 2.1833204236561405
epoch- 14/20, loss- 2.126685417249707
epoch- 15/20, loss- 2.074145279843294
epoch- 16/20, loss- 2.025608727264474
epoch- 17/20, loss- 1.9808416425968192
epoch- 18/20, loss- 1.9393367056019195
epoch- 19/20, loss- 1.9008196130734576
Accuracy on test set: 0.475

training for learning rate 0.0001

epoch- 0/20, loss- 3.1264896716609902
epoch- 1/20, loss- 2.4160236187294015
epoch- 2/20, loss- 2.035261817105016
epoch- 3/20, loss- 1.786417148535075
epoch- 4/20, loss- 1.6194850399339846
epoch- 5/20, loss- 1.4959622072431815
epoch- 6/20, loss- 1.393900048439698
epoch- 7/20, loss- 1.3195611691100428
epoch- 8/20, loss- 1.2649386407476704
epoch- 9/20, loss- 1.2209092057265276
epoch- 10/20, loss- 1.1845287611648831
epoch- 11/20, loss- 1.1568863417914081
epoch- 12/20, loss- 1.1351921381055659
epoch- 13/20, loss- 1.1164237794398042
epoch- 14/20, loss- 1.1000326378147598
epoch- 15/20, loss- 1.0844872668132723
epoch- 16/20, loss- 1.07091943404348
epoch- 17/20, loss- 1.0586150092167474
epoch- 18/20, loss- 1.0481481013913445
epoch- 19/20, loss- 1.0395281525356364
Accuracy on test set: 0.7065

training for learning rate 0.001




```
-----  
epoch- 0/20, loss- 1.9497445642901394  
epoch- 1/20, loss- 1.2993384825084473  
epoch- 2/20, loss- 1.2690101114737204  
epoch- 3/20, loss- 1.4184427706868088  
epoch- 4/20, loss- 1.7578005204110492  
epoch- 5/20, loss- 1.8468606362228392  
epoch- 6/20, loss- 1.9187014761793975  
epoch- 7/20, loss- 1.9301476383187912  
epoch- 8/20, loss- 1.954684354422394  
epoch- 9/20, loss- 2.0015675741810988  
epoch- 10/20, loss- 2.0789265015996974  
epoch- 11/20, loss- 2.1339791153075995  
epoch- 12/20, loss- 2.189122715648856  
epoch- 13/20, loss- 2.2454837146186017  
epoch- 14/20, loss- 2.3074667264063744  
epoch- 15/20, loss- 2.374467634050303  
epoch- 16/20, loss- 2.4610298487075  
epoch- 17/20, loss- 2.513137536375519  
epoch- 18/20, loss- 2.5639378346840944  
epoch- 19/20, loss- 2.6267847113492446  
Accuracy on test set: 0.4175
```


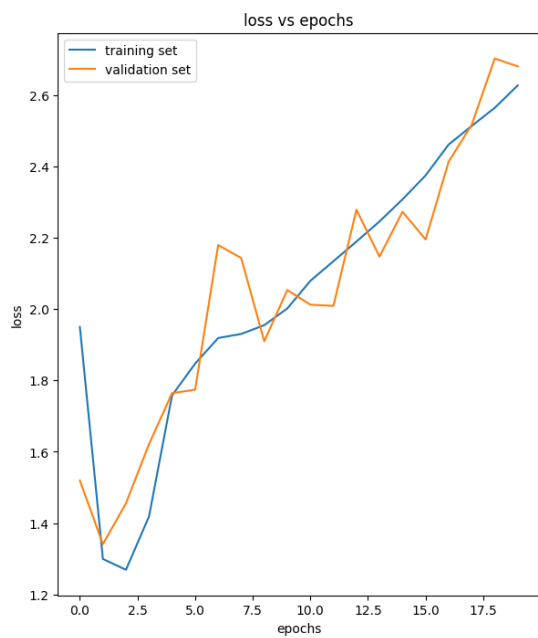


training for learning rate 0.01





```
-----  
epoch- 0/20, loss- 2.5788517836011904  
epoch- 1/20, loss- 2.2929408644279645  
epoch- 2/20, loss- 2.2326095191614663  
epoch- 3/20, loss- 2.1993326158445803  
epoch- 4/20, loss- 2.2087483499138463  
epoch- 5/20, loss- 2.230298142261838  
epoch- 6/20, loss- 2.218843048178565  
epoch- 7/20, loss- 2.214882221150696  
epoch- 8/20, loss- 2.2794222573451495  
epoch- 9/20, loss- 2.303701260303057  
epoch- 10/20, loss- 2.3083868867327473  
epoch- 11/20, loss- 2.360445903204188  
epoch- 12/20, loss- 2.315747718580962  
epoch- 13/20, loss- 2.352287978797347  
epoch- 14/20, loss- 2.351036021762887  
epoch- 15/20, loss- 2.3689706779446227  
epoch- 16/20, loss- 2.4134146906134775  
epoch- 17/20, loss- 2.4151051006850994  
epoch- 18/20, loss- 2.3850412245185364  
epoch- 19/20, loss- 2.3806079852477118  
Accuracy on test set: 0.565
```

training for learning rate 0.1

```
-----  
epoch- 0/20, loss- 38.90628817100049  
epoch- 1/20, loss- 3.4125694713350985  
epoch- 2/20, loss- 3.3861176871148966  
epoch- 3/20, loss- 3.3730124479257944  
epoch- 4/20, loss- 3.365337128801114  
epoch- 5/20, loss- 3.3604704919884516  
epoch- 6/20, loss- 3.3572651692551925  
epoch- 7/20, loss- 3.355133953504586  
epoch- 8/20, loss- 3.3537430802243335  
epoch- 9/20, loss- 3.352887837873292  
epoch- 10/20, loss- 3.3524351530959255  
epoch- 11/20, loss- 3.3522944384550475  
epoch- 12/20, loss- 3.352401659914012  
epoch- 13/20, loss- 3.3527101028588313  
epoch- 14/20, loss- 3.3531847573748714  
epoch- 15/20, loss- 3.3537987649954952  
epoch- 16/20, loss- 3.354531092717064  
epoch- 17/20, loss- 3.355364965636401  
epoch- 18/20, loss- 3.356286783956041  
epoch- 19/20, loss- 3.3572853580800017  
Accuracy on test set: 0.0315
```







6 for different epochs on tanh activation



```
[38]: # lets run it for 100 epochs so that we can observe what is happening for all
      ↪ the epochs between 20 and 100
      # neural network for tanh activation function
nn= MLP()
nn.add(fully_connected_layer(16, 40))
nn.add(activation_layer(act_func="tanh"))
#nn.add(activation_layer(activation_tanh, gradient_tanh))
nn.add(fully_connected_layer(40, 26))
nn.add(softmax_layer(26))
nn.use(cross_entropy_loss, grad_cross_entropy_loss)
nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=100,
      ↪ learning_rate=0.001, optimizer="adam")
nn.plot()

y_test_pred = np.array(nn.predict(x_test_arr))
y_test_pred = onehot.inverse_transform(y_test_pred)
y_test_true = onehot.inverse_transform(y_test_arr)
print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))
```

```
epoch- 0/100, loss- 2.5083383894457993
epoch- 1/100, loss- 1.862494219274102
epoch- 2/100, loss- 1.859269353058428
epoch- 3/100, loss- 1.7897892848493262
epoch- 4/100, loss- 1.6670168525369806
epoch- 5/100, loss- 1.5671354038545084
epoch- 6/100, loss- 1.504174246043134
epoch- 7/100, loss- 1.4618135603814795
epoch- 8/100, loss- 1.434236305329016
epoch- 9/100, loss- 1.4185731151648417
epoch- 10/100, loss- 1.4088692563745076
epoch- 11/100, loss- 1.4041577427653387
epoch- 12/100, loss- 1.3947059373933794
epoch- 13/100, loss- 1.3630910813052266
epoch- 14/100, loss- 1.3534617727820148
epoch- 15/100, loss- 1.3342925821242186
epoch- 16/100, loss- 1.3260759747736384
epoch- 17/100, loss- 1.3207539992364896
epoch- 18/100, loss- 1.3125294417135311
epoch- 19/100, loss- 1.312032116817627
epoch- 20/100, loss- 1.3214866073243292
epoch- 21/100, loss- 1.315682243804574
epoch- 22/100, loss- 1.3145655307803583
epoch- 23/100, loss- 1.3192780988413066
epoch- 24/100, loss- 1.3299791663405724
epoch- 25/100, loss- 1.3319387623831678
epoch- 26/100, loss- 1.357138070262373
```

epoch- 27/100, loss- 1.3504401593634185
epoch- 28/100, loss- 1.3558825960137624
epoch- 29/100, loss- 1.3549569719925176
epoch- 30/100, loss- 1.3594377434534792
epoch- 31/100, loss- 1.3684717927379018
epoch- 32/100, loss- 1.3657308543948155
epoch- 33/100, loss- 1.3844536920074386
epoch- 34/100, loss- 1.3914836814672764
epoch- 35/100, loss- 1.40782351115854
epoch- 36/100, loss- 1.404206684369323
epoch- 37/100, loss- 1.423957260504462
epoch- 38/100, loss- 1.4357782687157392
epoch- 39/100, loss- 1.4365907805142784
epoch- 40/100, loss- 1.445290594323743
epoch- 41/100, loss- 1.4429183119882876
epoch- 42/100, loss- 1.4696953877383216
epoch- 43/100, loss- 1.4653138667510746
epoch- 44/100, loss- 1.4690922553783947
epoch- 45/100, loss- 1.4746835543125512
epoch- 46/100, loss- 1.4686291973406667
epoch- 47/100, loss- 1.4892202837996047
epoch- 48/100, loss- 1.4976792152769425
epoch- 49/100, loss- 1.4948206664702868
epoch- 50/100, loss- 1.508604169954548
epoch- 51/100, loss- 1.5264147681023776
epoch- 52/100, loss- 1.5603928761102697
epoch- 53/100, loss- 1.542329678881554
epoch- 54/100, loss- 1.5599252949773672
epoch- 55/100, loss- 1.5906956581534368
epoch- 56/100, loss- 1.5722446125534717
epoch- 57/100, loss- 1.5865526798660587
epoch- 58/100, loss- 1.5599227160194389
epoch- 59/100, loss- 1.5714872162176285
epoch- 60/100, loss- 1.601489532084466
epoch- 61/100, loss- 1.6243578376099534
epoch- 62/100, loss- 1.631387137451445
epoch- 63/100, loss- 1.6123038017073656
epoch- 64/100, loss- 1.6220277429824603
epoch- 65/100, loss- 1.6140262343176162
epoch- 66/100, loss- 1.6267251630054298
epoch- 67/100, loss- 1.599705748892451
epoch- 68/100, loss- 1.5917657768223648
epoch- 69/100, loss- 1.5994445826426633
epoch- 70/100, loss- 1.5835260924243775
epoch- 71/100, loss- 1.5587780400876017
epoch- 72/100, loss- 1.5862024200151421
epoch- 73/100, loss- 1.5813871261313568
epoch- 74/100, loss- 1.5734797259767828

epoch- 75/100, loss- 1.587176553318264
epoch- 76/100, loss- 1.6078836383998318
epoch- 77/100, loss- 1.6145763727587612
epoch- 78/100, loss- 1.6003365976903736
epoch- 79/100, loss- 1.6169109294116213
epoch- 80/100, loss- 1.5812456187430421
epoch- 81/100, loss- 1.6001879995341592
epoch- 82/100, loss- 1.5776605381431168
epoch- 83/100, loss- 1.575021584627835
epoch- 84/100, loss- 1.6080562780049528
epoch- 85/100, loss- 1.5776439534206919
epoch- 86/100, loss- 1.569789429635737
epoch- 87/100, loss- 1.5922797398132431
epoch- 88/100, loss- 1.572955349533713
epoch- 89/100, loss- 1.5860103455637136
epoch- 90/100, loss- 1.573055343202307
epoch- 91/100, loss- 1.5901317812996927
epoch- 92/100, loss- 1.6230099808444323
epoch- 93/100, loss- 1.6084203990571568
epoch- 94/100, loss- 1.6294059000454038
epoch- 95/100, loss- 1.6446473757210134
epoch- 96/100, loss- 1.6373783759819138
epoch- 97/100, loss- 1.6552147859877746
epoch- 98/100, loss- 1.6621141059444748
epoch- 99/100, loss- 1.6611453152746687
Accuracy on test set: 0.683



7 for different epochs on relu activation

```
[39]: # lets run it for 100 epochs so that we can observe what is happening for all
      ↪ the epochs between 20 and 100
      # neural network for tanh activation function
nn= MLP()
nn.add(fully_connected_layer(16, 40))
nn.add(activation_layer(act_func="relu"))
#nn.add(activation_layer(activation_tanh, gradient_tanh))
nn.add(fully_connected_layer(40, 26))
nn.add(softmax_layer(26))
nn.use(cross_entropy_loss, grad_cross_entropy_loss)
nn.fit(x_train_arr, y_train_arr, x_valid_arr, y_valid_arr, epochs=100,
      ↪ learning_rate=1e-4, optimizer="adam")
nn.plot()

y_test_pred = np.array(nn.predict(x_test_arr))
y_test_pred = onehot.inverse_transform(y_test_pred)
y_test_true = onehot.inverse_transform(y_test_arr)
print("Accuracy on test set: ", accuracy_score(y_test_pred,y_test_true))
```


```
epoch- 0/100, loss- 2.9884792852037654
epoch- 1/100, loss- 2.113862652356504
epoch- 2/100, loss- 1.759463925354568
epoch- 3/100, loss- 1.5623163976989092
epoch- 4/100, loss- 1.441973814711572
epoch- 5/100, loss- 1.3627454071201852
epoch- 6/100, loss- 1.3048068553075987
epoch- 7/100, loss- 1.2620756334812169
epoch- 8/100, loss- 1.2307569469413993
epoch- 9/100, loss- 1.2053166638904065
epoch- 10/100, loss- 1.1829633575528484
epoch- 11/100, loss- 1.1636753929978416
epoch- 12/100, loss- 1.1457332535655758
epoch- 13/100, loss- 1.1271035167026071
epoch- 14/100, loss- 1.1085508052583206
epoch- 15/100, loss- 1.0906723336778799
epoch- 16/100, loss- 1.0767276665691907
epoch- 17/100, loss- 1.070428385252438
epoch- 18/100, loss- 1.0703047058086363
epoch- 19/100, loss- 1.0737989161247021
epoch- 20/100, loss- 1.0818209878601366
epoch- 21/100, loss- 1.1012466393637723
epoch- 22/100, loss- 1.1325516638786413
epoch- 23/100, loss- 1.1778724890545886
epoch- 24/100, loss- 1.2256920166974854
epoch- 25/100, loss- 1.2634608351492673
epoch- 26/100, loss- 1.2967258859045598
```


epoch- 27/100, loss- 1.346000912051773
epoch- 28/100, loss- 1.3815850281039908
epoch- 29/100, loss- 1.4179833197163512
epoch- 30/100, loss- 1.465804479055009
epoch- 31/100, loss- 1.5067224531453733
epoch- 32/100, loss- 1.544306507465588
epoch- 33/100, loss- 1.5671798563167443
epoch- 34/100, loss- 1.5945198128103195
epoch- 35/100, loss- 1.5969278468485326
epoch- 36/100, loss- 1.5924839259890582
epoch- 37/100, loss- 1.5796826742605432
epoch- 38/100, loss- 1.5653434004237547
epoch- 39/100, loss- 1.5423338320589415
epoch- 40/100, loss- 1.5709006630753735
epoch- 41/100, loss- 1.629656109995847
epoch- 42/100, loss- 1.6514606375513885
epoch- 43/100, loss- 1.6817803377440712
epoch- 44/100, loss- 1.7098601818282364
epoch- 45/100, loss- 1.679425483283555
epoch- 46/100, loss- 1.6614759210784895
epoch- 47/100, loss- 1.6591222536192698
epoch- 48/100, loss- 1.665200350769158
epoch- 49/100, loss- 1.6431206317437204
epoch- 50/100, loss- 1.6546449922734967
epoch- 51/100, loss- 1.6509376891435212
epoch- 52/100, loss- 1.6522502761273388
epoch- 53/100, loss- 1.6587093281261787
epoch- 54/100, loss- 1.6644064191536572
epoch- 55/100, loss- 1.6833214480682657
epoch- 56/100, loss- 1.684075892386601
epoch- 57/100, loss- 1.6726669313748885
epoch- 58/100, loss- 1.6729470677838698
epoch- 59/100, loss- 1.6942009722376274
epoch- 60/100, loss- 1.6880536497880307
epoch- 61/100, loss- 1.6819840792617122
epoch- 62/100, loss- 1.6930929856515966
epoch- 63/100, loss- 1.704510889502957
epoch- 64/100, loss- 1.7571859829430043
epoch- 65/100, loss- 1.7672116706523893
epoch- 66/100, loss- 1.7572427424386927
epoch- 67/100, loss- 1.756441447464374
epoch- 68/100, loss- 1.7742967326961234
epoch- 69/100, loss- 1.776524370092036
epoch- 70/100, loss- 1.79617753228131
epoch- 71/100, loss- 1.8055971715976997
epoch- 72/100, loss- 1.8229772687416854
epoch- 73/100, loss- 1.8267710746031556
epoch- 74/100, loss- 1.845285428619157

```

epoch- 75/100, loss- 1.8450792971857672
epoch- 76/100, loss- 1.8763156725530323
epoch- 77/100, loss- 1.8812410202838281
epoch- 78/100, loss- 1.9089200764353194
epoch- 79/100, loss- 1.9026066778792041
epoch- 80/100, loss- 1.9187547021535636
epoch- 81/100, loss- 1.9203238853444295
epoch- 82/100, loss- 1.9399640845917359
epoch- 83/100, loss- 1.9199544543986087
epoch- 84/100, loss- 1.9201956225087153
epoch- 85/100, loss- 1.9184062158669166
epoch- 86/100, loss- 1.9180194939988522
epoch- 87/100, loss- 1.9391101290660304
epoch- 88/100, loss- 1.9232234762458822
epoch- 89/100, loss- 1.9482137079016997
epoch- 90/100, loss- 1.9292564630366065
epoch- 91/100, loss- 1.9385419097693577
epoch- 92/100, loss- 1.9616793154322212
epoch- 93/100, loss- 1.9567109167194676
epoch- 94/100, loss- 1.9459618471889641
epoch- 95/100, loss- 1.9448487839507769
epoch- 96/100, loss- 1.9410503508121828
epoch- 97/100, loss- 1.9411272654817069
epoch- 98/100, loss- 1.9321444982518572
epoch- 99/100, loss- 1.9356582689177622
Accuracy on test set: 0.519

```



so from the above analysis we ge that the best parameters for tanh are as follows- 1. learning rate- 0.001 or $1e-3$ 2. epoch- 100 epochs can be taken as test accuracy and generalization gap does not

change significantly we get accuracy on test set- 68.31. learning rate- 0.0001 or $1e-4$ 2. epoch- we must take 20 epochs because after that the loss increases and the accuracy decreases rapidly we get accuracy on test set- 63.5

8 finally we prefer tanh activation with learning rate 0.001 or $1e-3$, number of epochs 100 and get test accuracy of 68.3%

```
[40]: # code ended for Q.1.
```